

Deep Learning in Python (An Example in NLP)

Yuan Luo

Assistant Professor

Department of Preventive Medicine

Departments of IEMS and EECS (Courtesy)

Northwestern University

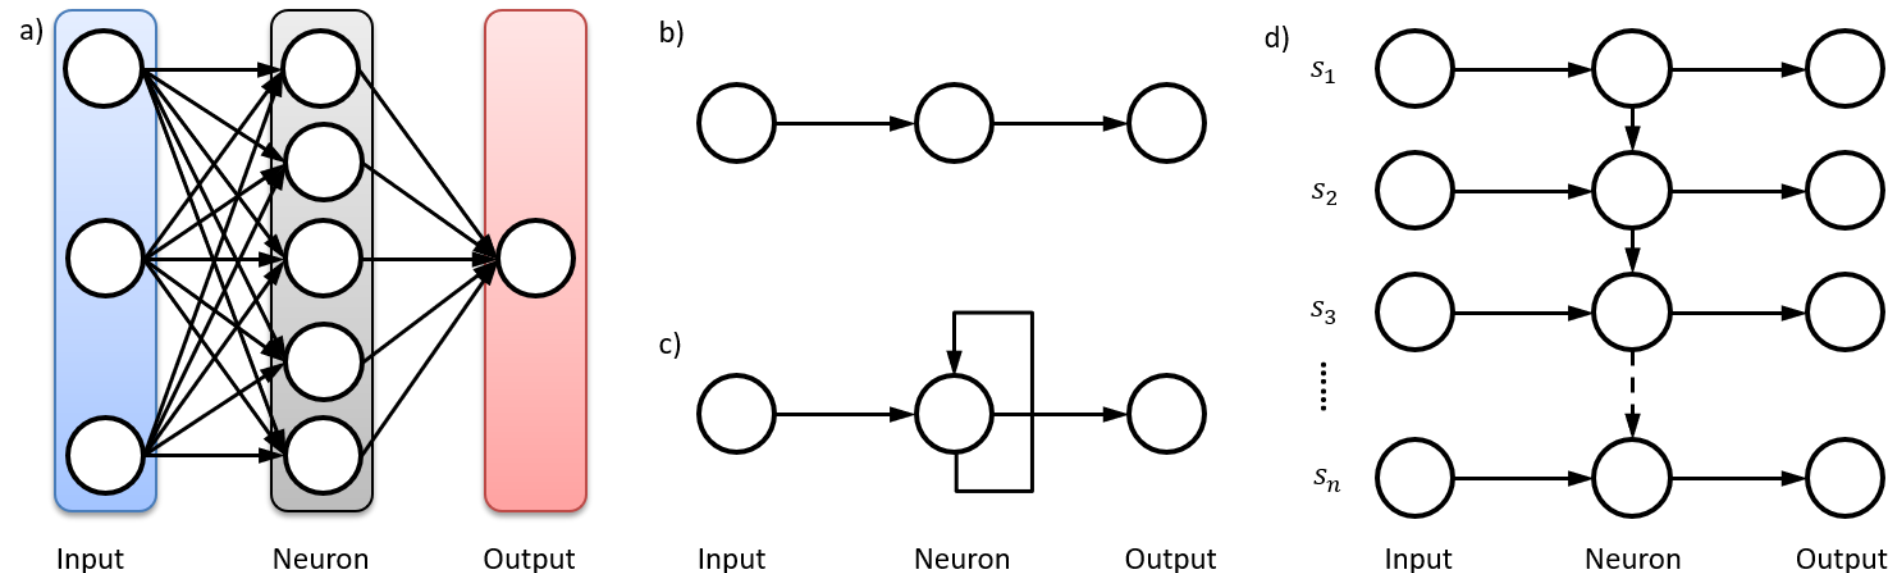
yuan.luo@northwestern.edu

 @yuanhypnosluo


© Yuan Luo. All Rights Reserved

Basics on Neural Networks

- From multi-layer neural networks to recurrent neural networks



Convolutional Neural Networks



Y Kim. Convolutional neural networks for sentence classification. *In EMNLP 2014*, 1746-1751

Using GCN for Long Text Understanding

- Challenges of Long Text Understanding
 - CNN and RNN prioritize locality and sequentiality.
 - They can model local consecutive word sequences well
 - They may ignore global word co-occurrence in a corpus
- GCN can
 - Generalizing well-established neural network models like CNN that apply to regular grid structure (2-d mesh or 1-d sequence) to work on arbitrarily structured graphs
 - Can preserve global structure information of a graph in graph embeddings (node, edge, subgraph and whole graph embeddings)

Graph Convolution

- Denote $A \in \mathbb{R}^{n \times n}$ as the dependency graph G 's adjacency matrix
- Graph Laplacian is $L = D - A$, where $D_{ii} = \sum_j A_{ij}$
- Let $U \in \mathbb{R}^{n \times n}$ be the matrix of eigenvectors of the normalized graph Laplacian $L_n = I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}} = U\Lambda U^T$
- Let $g_w \in \mathbb{R}^{n \times n}$ be a **Fourier domain** filter matrix parametrized with a scalar w as its diagonal elements (**recall signal processing theory**)
- The graph convolution for the 1-dimensional embedding $x \in \mathbb{R}^n$ (for n words) is
$$h = U g_w U^T x = U \text{diag}([w, \dots, w]) U^T x = U U^T x \text{diag}([w, \dots, w])$$
- Simplify using Chebyshev polynomial approximation
- $h = (I + D^{-\frac{1}{2}}AD^{-\frac{1}{2}})xw$, after renormalization $h = \tilde{D}^{-\frac{1}{2}}\tilde{A}\tilde{D}^{-\frac{1}{2}}xw$
- Extend the embedding and convolved signal to d -dimensional $X \in \mathbb{R}^{n \times d}$ and $H \in \mathbb{R}^{n \times d}$

$$H = \rho(\tilde{D}^{-\frac{1}{2}}\tilde{A}\tilde{D}^{-\frac{1}{2}}XW)$$

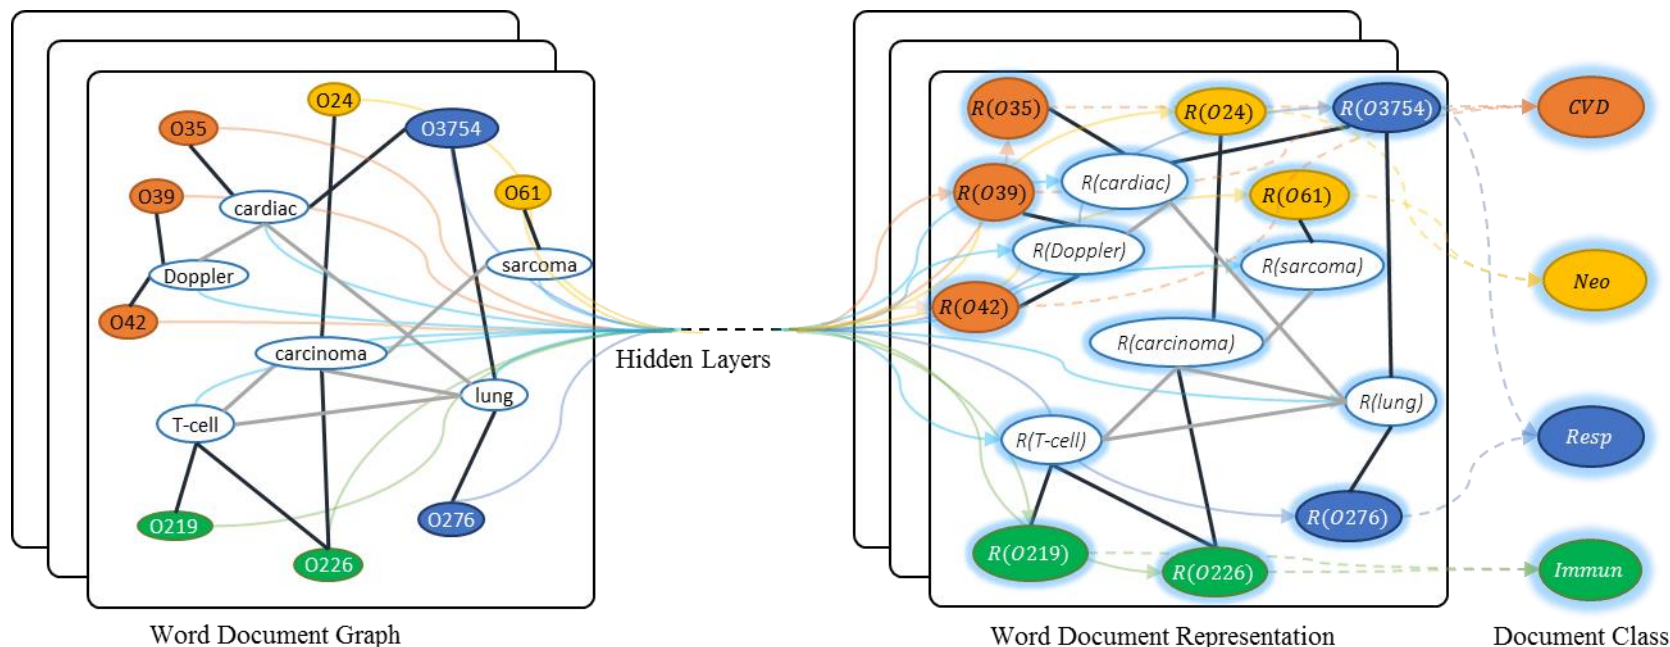
Graph Convolutional Networks (GCN)

- A graph $G = (V, E)$:
 - $(v, v) \in E$ for any v
 - $X \in R^{n \times m}$: node features matrix
 - A : adjacency matrix, degree matrix $D_{ii} = \sum_j A_{jj}$
 - $\tilde{A} = \tilde{D}^{-\frac{1}{2}} A \tilde{D}^{-\frac{1}{2}}$: normalized symmetric adjacency matrix
 - W_j : weight matrix, trained via SGD
- One layer GCN:
- $L^{(1)} = \rho(\tilde{A}XW_0)$
- Stacking multiple GCN layers:
- $L^{(j+1)} = \rho(\tilde{A}L^{(j)}W_j)$

Graph Convolutional Networks (GCN)

- GCN can capture information only about immediate neighbors with one layer
- When multiple GCN layers are stacked, one can incorporate higher order neighborhoods information
 - e.g., a two-layers GCN can allow message passing among nodes that are at maximum two steps away.
- A special form of Laplacian smoothing:
 - computes the new features of a node as the weighted average of itself and its neighbors (second order neighbors for a two-layer GCN).

Text Graph Convolutional Networks (Text GCN)



L Yao, C Mao, **Y Luo***. Graph Convolutional Networks for Text Classification. *Proceedings of AAAI Conference on Artificial Intelligence 2019 Full paper*.

Text Graph Convolutional Networks (Text GCN)

- Document content and global word co-occurrence
 - Document-word edges: TF-IDF
 - Word-word edges: point-wise mutual information (PMI)

$$p(i, j) = \frac{\#W(i, j)}{\#W}$$
$$p(i) = \frac{\#W(i)}{\#W}$$
$$\text{PMI}(i, j) = \log \frac{p(i, j)}{p(i)p(j)}$$

$$A_{ij} = \begin{cases} \text{PMI}(i, j) & i, j \text{ are words, } \text{PMI}(i, j) > 0 \\ \text{TF-IDF}_{ij} & i \text{ is document, } j \text{ is word} \\ 1 & i = j \\ 0 & \text{otherwise} \end{cases}$$

Text Graph Convolutional Networks (Text GCN)

- A simple two-layer GCN:
 - one-hot feature matrix for words and documents: $X = I$
 - 1st layer document and word embeddings: $\tilde{A}XW_0$
 - 2nd layer document and word embeddings: $\tilde{A}\text{ReLU}(\tilde{A}XW_0)W_1$
 - \mathcal{Y}_D is the set of document indices that have labels and F is the dimension of the output features, which is equal to the number of classes, Y is the label indicator matrix
 - Loss function

$$Z = \text{softmax}(\tilde{A} \text{ReLU}(\tilde{A}XW_0)W_1)$$

$$\mathcal{L} = - \sum_{d \in \mathcal{Y}_D} \sum_{f=1}^F Y_{df} \ln Z_{df}$$

Dataset

- The Ohsumed corpus is from the MEDLINE database, which is a bibliographic database of important medical literature maintained by the National Library of Medicine
- In this tutorial, we created a subsample of the 2,762 unique diseases abstracts from 3 categories
 - C04: Neoplasms
 - C10: Nervous System Diseases
 - C14: Cardiovascular Diseases
- As we focus on single-label text classification, the documents belonging to multiple categories are excluded
- 1230 train (use 10% as validation), 1532 test

Now let's look at some code

- https://github.com/yuanluo/text_gcn_tutorial
- `Run python remove_words.py ohsumed_3`
- `Run python build_graph.py ohsumed_3`
- `Run python train.py ohsumed_3`
- `Run python tsne.py`

Results

- C04: Neoplasms
- C10: Nervous System Diseases
- C14: Cardiovascular Diseases

